

# Parallel Debugging

Jonathan Schmalfuß

Chair of Scientific Computing  
University of Bayreuth

March 19, 2025

“Sequential programming is really hard, and parallel programming is a step beyond that.” - Andrew S. Tanenbaum, professor at Vrije Universiteit Amsterdam

“Debugging is twice as hard as writing the code in the first place. Therefore, if you write the code as cleverly as possible, you are, by definition, not smart enough to debug it.” - Brian Kernighan, professor at Princeton University.

“Most bugs also appear in the sequential version of the code” - me

“Most bugs also appear in the sequential version of the code” - me

## Distinguish Problem

Sequential  
Problem

←  
problem  
persists

Force Sequential  
Execution via  
`MPI_Barrier()`

→  
works  
now

Parallel  
Problem

“Most bugs also appear in the sequential version of the code” - me

## Distinguish Problem

Sequential  
Problem

←  
problem  
persists

Force Sequential  
Execution via  
`MPI_Barrier()`

→  
works  
now

Parallel  
Problem

The experts opinion: Anthony Williams - author/ coauthor of the thread library in C++

- 1 Reviewing code to locate potential bugs
- 2 Locating concurrency-related bugs by testing / Designing for testability

- Are there any ordering requirements between the operations done in this process and those done in another? How are those requirements enforced?
- Which data needs to be protected from concurrent access? How do you ensure that the data is protected?
- Where in the code could other processes be at this time?
- Is the data loaded by this process still valid?
- If you assume that another process could be modifying the data, what would that mean and how could you ensure that this never happens?

# Parallel Debugging: How to?

## the easy way

- use a debugger and or fronted made for debugging MPI code
- Industry Standard: `ddt`<sup>1</sup>, `TotalView`<sup>1</sup>
- Open Source Projects / Free: `mdb`, oneAPI with `mpigdb` or a shell script `tmapi`

## should-always-work way

- minimal requirements: a debugger (gdb) + a way of finding the running processes (ps/top)
- `mpirun` creates multiple processes → attach to relevant processes or all → debug each of them sequentially
- limits: low number of processes, requires duplicating input for each process

<sup>1</sup>temporary free / student license available

## Debug Deadlocks: Attach to the process [Live-session]

Situation: you have a deadlock, i.e. your executable is stuck

- 1 Compile with debug flags: `mpic++ -g -Wall <file> -o <name>`
- 2 Wait until stuck
- 3 Figure out process id's via `top` or `ps -a | grep <name>`
- 4 Attach to the process and see where you are stuck → figure out what the problem is



## Debug Deadlocks: Attach to the process [Live-session]

Situation: you have a deadlock, i.e. your executable is stuck

- 1 Compile with debug flags: `mpic++ -g -Wall <file> -o <name>`
- 2 Wait until stuck
- 3 Figure out process id's via `top` or `ps -a | grep <name>`
- 4 Attach to the process and see where you are stuck → figure out what the problem is

## Attaching debugger to several instances of the executable

- Use `mpirun` to launch separate instances of serial debuggers
- Drawback: many process, usually problematic

OpenMPI: FAQ: Debugging applications in parallel

- Requires: Open MPI 1.3 or later, and Valgrind 3.2.0 or later
- Otherwise: works, but with many false positives
- Needs to be enabled at compilation state of OpenMPI, unfortunately often is not
- to enable locally, see [How can I use Memchecker](#)
- `mpirun -np 2 valgrind`
  - ↳ `--suppressions=$PREFIX/share/openmpi/openmpi-valgrind.supp`
  - ↳ `<executable name>`

ddt [Live-session]

- same deadlock problem

Brought you some programs, which you can check out. Find the errors.

- 1 Is it a sequential problem or parallel for `mpi_reusing_a_buffer.cpp`